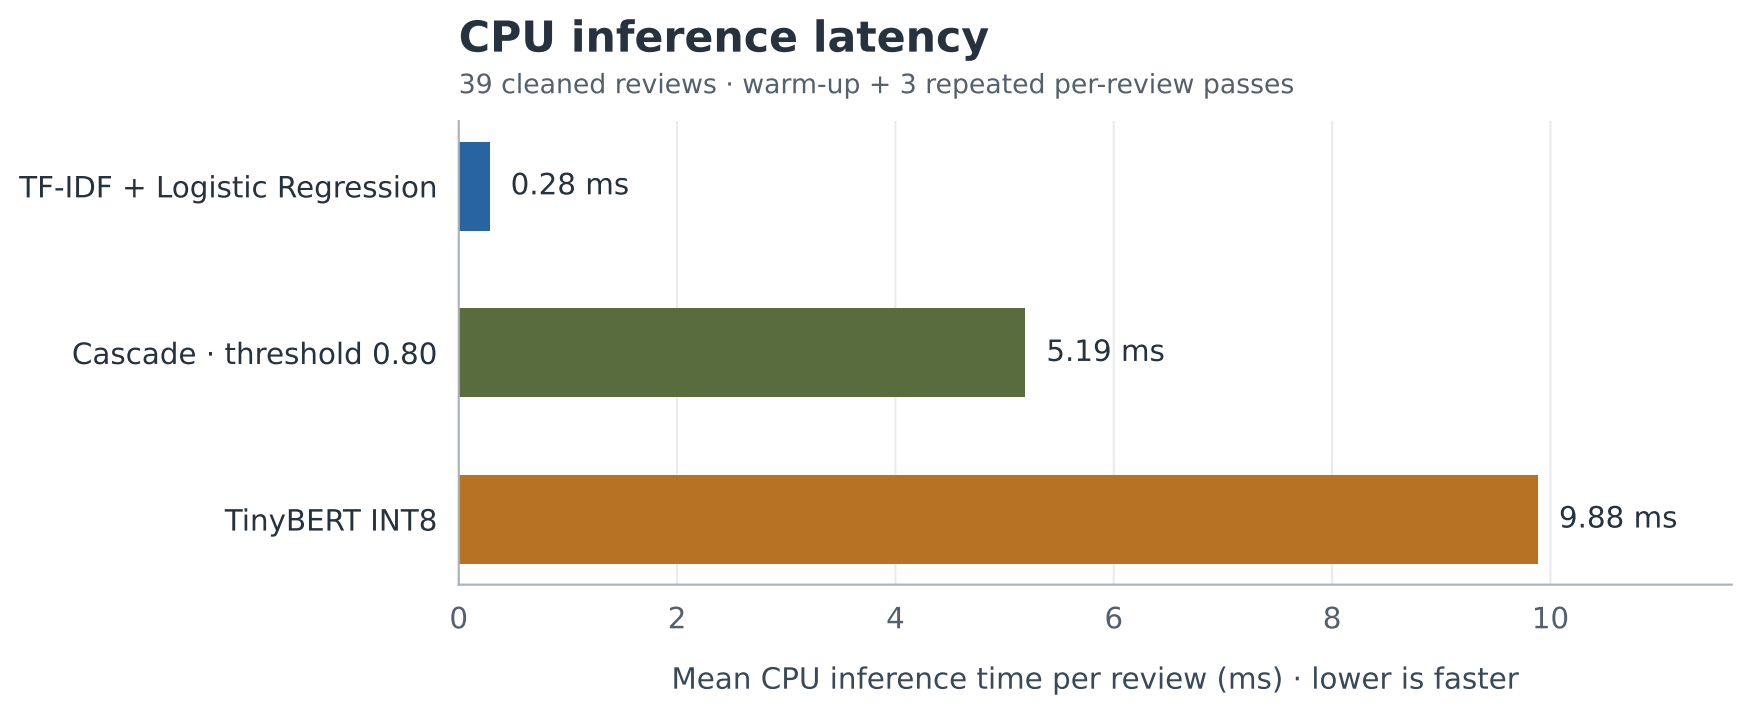
CPU inference latency
39 cleaned reviews · warm-up + 3 repeated per-review passes
TF-IDF + Logistic Regression
0.28 ms
Cascade · threshold 0.80
5.19 ms
TinyBERT INT8
9.88 ms
0
2
4
6
8
10
Mean CPU inference time per review (ms) · lower is faster